



Data Science Intern at Data Glacier

Project: Hate Speech Detection using Transformers (Deep Learning)

Week 10: Deliverable

Group Name: ToxiTrackers

Name: Mehedi Hasan Shakil

University: Chittagong University of Engineering & Technology

Email: shakilcuetcse@gmail.com

Country: Bangladesh

Specialization: Data Science

Batch Code: LISUM44

Date: 09 June 2025

Submitted to: Data Glacier

Table of Contents:

1. Project Plan	3
2. Problem Description	3
3. Business Understanding	3
4. Data Preprocessing.....	4
4.1 Data Cleaning.....	4
4.2 Remove Stopwords.....	5
4.3 Tokenization	5
4.4 Lemmatization.....	5
5. Feature Extraction.....	5
5.1 Bag of Words(BoW).....	6
5.2 TF-IDF	6
5.3 Word Embeddings.....	6
5.4 BERT.....	6
6. EDA (Exploratory Data Analysis).....	6
6.1 Number of Words per Tweet.....	7
6.2 Most Frequent Words in the Entire Dataset.....	7
6.3 Most Frequent Words in Hate vs Non-Hate Tweets.....	8
6.4 Word Cloud for Each Class.....	9
6.5 Hate vs. Non-Hate Tweet Distribution.....	9
6.6 Handling Class Imbalance.....	10

1. Project Plan

Weeks	Date	plan
Weeks 07	May 19, 2025	Problem Description, Business understanding, Data Intake Report
Weeks 08	May 26, 2025	Data Preprocessing
Weeks 09	June 2, 2025	Data Cleansing and Transformation
Weeks 10	June 9, 2025	EDA performed
Weeks 11	June 16, 2025	EDA Presentation and Proposed Modelling Technique
Weeks 12	June 23, 2025	Model Selection and Model Building
Weeks 13	June 30, 2025	Final Submission (Report + Code + Presentation)

2. Problem Description

Hate speech poses a significant challenge on digital platforms, often manifesting as verbal, written, or behavioral communication targeting individuals or groups based on religion, ethnicity, nationality, race, gender, or other identity traits. With the rise of social media usage, particularly platforms like Twitter, monitoring and mitigating such harmful content has become essential.

This project aims to build an automated hate speech detection system using advanced Natural Language Processing (NLP) techniques and transformer-based deep learning models. The task is framed as a binary text classification problem, where the model learns to distinguish between hateful and non-hateful tweets. The goal is to create a reliable and scalable model capable of supporting content moderation efforts by detecting hate speech in real time.

3. Business Understanding

Social media platforms are essential tools for global communication but also serve as mediums for the spread of hate speech and abusive content. The unchecked proliferation of such content can lead to social unrest, legal issues, and reputational damage for platform providers.

Businesses operating these platforms are increasingly under pressure to ensure user safety and compliance with regulations on digital conduct.

By deploying an automated hate speech detection system, companies can significantly enhance their moderation capabilities. This not only helps maintain a safer online environment but also improves user trust and brand integrity. From a data science perspective, this project addresses key business goals such as minimizing manual moderation efforts, reducing harmful interactions, and supporting ethical AI practices in content governance.

4. Data Preprocessing

In this section, we describe the text preprocessing techniques applied to prepare the dataset for machine learning models. Raw text data typically contains various inconsistencies, noise, and irrelevant components that can negatively affect the performance of NLP-based models. Therefore, a sequence of cleaning steps was carried out to normalize and refine the input tweets.

4.1 Text Cleaning

Text data collected from social media platforms such as Twitter often contains unwanted elements like user mentions, hyperlinks, emojis, and non-standard formatting. Cleaning such data is essential for ensuring better feature extraction and model performance. The steps below summarize our cleaning pipeline:

4.1.1 Lowercasing

All characters in the text were converted to lowercase. This normalization step ensures that words like “*Hate*” and “*hate*” are treated identically. Without this transformation, these words would be considered different features in the model, leading to unnecessary sparsity and dimensionality in the feature space.

4.1.2 Removing User Mentions

Twitter handles or user mentions (e.g., *@username*) were removed. These references are unique to individual users and do not contribute semantically to the classification task. Their removal helps avoid model distraction from non-informative tokens.

4.1.3 Removing URLs

URLs (web links) were stripped from the text. Since the task focuses on analyzing the sentiment or intent of the written content, web addresses—which usually lead to external content—do not provide meaningful context and are treated as noise.

4.1.4 Removing Special Characters and Digits

Special characters (e.g., @, #, !, %) and numeric digits were eliminated. These characters typically do not carry linguistic meaning in the context of sentiment or hate speech detection. A regular expression-based filter was applied to retain only alphabetic characters and necessary whitespace.

4.1.5 Stripping Whitespace

Any leading or trailing spaces in the tweets were removed. This ensures that the text is properly formatted and avoids inconsistencies during tokenization and further processing.

4.2 Remove Stopwords

Stopwords are common words such as “the”, “is”, “in”, “and”, etc., which do not carry significant meaning and are often removed to reduce noise in the text. In this step, the stopwords were removed from each tweet using the Natural Language Toolkit (NLTK)'s built-in English stopword list. This step helps to focus the analysis on the most relevant and informative words.

4.3 Tokenization

Tokenization involves breaking down each tweet into individual tokens or words. This process is essential for further natural language processing (NLP) tasks such as lemmatization and vectorization. Tokenization was performed using NLTK's word tokenizer, which splits the cleaned text into word-level tokens while handling punctuation and whitespace effectively.

4.4 Lemmatization

Lemmatization is the process of reducing words to their base or dictionary form, known as a lemma. Unlike stemming, lemmatization considers the context of the word to ensure that it is grammatically correct. For instance, words like "running", "ran", and "runs" are reduced to the base form "run". This was performed using spaCy's lemmatizer, helping to normalize the text and reduce feature dimensionality.

5. Feature Extraction

After preprocessing the text data, several feature extraction techniques were employed to convert the textual data into numerical representations suitable for machine learning models. These representations capture different aspects of the text and were used as input features for classification.

5.1 Bag of Words (BoW)

The Bag of Words model is a simple and commonly used text representation technique. It creates a vocabulary of all unique words in the dataset and represents each tweet as a vector containing the frequency of each word. The CountVectorizer from scikit-learn was used to transform the preprocessed tweets into sparse vectors. This method captures the presence and frequency of words but ignores grammar and word order.

5.2 Term Frequency-Inverse Document Frequency (TF-IDF)

TF-IDF improves upon the Bag of Words model by not only considering the frequency of words but also their uniqueness across all documents. Words that appear frequently in one document but rarely across others are given higher weights. The TfidfVectorizer from scikit-learn was applied to the preprocessed text data to generate a TF-IDF matrix. This representation helps reduce the impact of common yet less informative words, making the feature space more meaningful for classification tasks.

5.3 Word Embeddings

To capture semantic relationships and contextual meanings, pre-trained word embeddings were used. Specifically, GloVe (Global Vectors for Word Representation) embeddings were loaded, and each tweet was represented as the average of its word vectors. This method enables the model to understand relationships between similar words (e.g., "king" and "queen") and provides dense, real-valued feature vectors that are more informative than BoW or TF-IDF.

5.4 BERT Embeddings

Bidirectional Encoder Representations from Transformers (BERT) is a deep learning model developed by Google that provides context-aware word embeddings. The bert-base-uncased pre-trained model was used via the Hugging Face Transformers library to extract embeddings. Each tweet was tokenized using BertTokenizer, and the [CLS] token's embedding was extracted from BertModel to represent the entire tweet. These embeddings are highly contextual and provide rich information suitable for downstream classification tasks. The extracted features can also be fine-tuned later with BERT's classification head for improved performance.

6. EDA (Exploratory Data Analysis)

In this part, we do a simple analysis to understand the data's structure and patterns. We check the number of total words in tweets and find the most frequent words appearing in hate and non-hate tweets. Then, we check the class distribution and use a resampling technique to solve the imbalance issue.

6.1 Number of Words per Tweet

We calculated the number of words in each tweet after the preprocessing steps. The distribution is visualized in *Figure 6.1*, which shows that the majority of tweets contain between 5 to 15 words. This indicates a relatively concise structure, which is typical for Twitter posts.




Figure 6.1: Histogram showing the distribution of word counts per tweet.

6.2 Most Frequent Words in the Entire Dataset

The 20 most frequently occurring words across the entire dataset are shown in Figure 6.2. These terms were extracted from the lemmatized token lists and represent the most commonly used language elements in both hate and non-hate tweets.




Figure 6.2: Bar plot of the top 20 most frequent words across all tweets.

6.3 Most Frequent Words in Hate vs Non-Hate Tweets

To distinguish the linguistic patterns between hate and non-hate tweets, we separately analyzed the most frequent words in each class. As shown in Figure 6.3, hate tweets tend to use more aggressive or targeted language, while non-hate tweets exhibit more neutral or conversational vocabulary.





Figure 6.3: Horizontal bar plots of the top 15 frequent words in hate vs. non-hate tweets.

Word clouds offer a visually intuitive way to understand the most prominent words in each category. Figure 6.4 shows the word clouds for hate and non-hate tweets, respectively. The size of each word in the cloud corresponds to its frequency. Hate-related tweets reveal more emotionally charged or offensive language, whereas non-hate tweets reflect diverse, less targeted content.



The original class distribution is visualized in Figure 6.5. It is evident that the dataset is imbalanced, with a significantly higher number of non-hate tweets compared to hate tweets. This imbalance can affect the model's performance if not properly addressed.



6.6 Handling Class Imbalance

To mitigate the impact of class imbalance, we applied upsampling to the minority class (hate tweets). After resampling, both classes have equal representation in the dataset, as shown in Figure 6.6. This step ensures that the model is trained on a balanced dataset, which improves its ability to generalize across both classes.




Figure 6.6: Class distribution after upsampling the minority class (hate tweets).

This EDA helped us better understand the nature of the dataset, uncover class imbalance, and prepare for effective model training in the subsequent steps.